

TEAM 16

Aravind Reddy Sheru

Email: asheru@iu.edu (<mailto:asheru@iu.edu>)

Sai Charan Chintala

Email: sachin@iu.edu (<mailto:sachin@iu.edu>)

Seongbo Sim

Email: simseo@iu.edu (<mailto:simseo@iu.edu>)

Yun Joo An

Email: yunjooan@iu.edu (<mailto:yunjooan@iu.edu>)

```

In [1]: import warnings
warnings.simplefilter('ignore')

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import re
from time import time
from scipy import stats
import json

from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.metrics import make_scorer, roc_auc_score, log_loss, accuracy_score
from sklearn.preprocessing import LabelEncoder

from sklearn.metrics import confusion_matrix
from IPython.display import display, Math, Latex

from sklearn.linear_model import Lasso, Ridge, LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb

# Read data from application_train dataset.
data = pd.read_csv('application_train.csv')
data.head()

```

```

Out[1]:

```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_O
0	100002	1	Cash loans	M	N	
1	100003	0	Cash loans	F	N	

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_O
2	100004	0	Revolving loans	M	Y	
3	100006	0	Cash loans	F	N	
4	100007	0	Cash loans	M	N	

5 rows × 122 columns

```
In [2]: y = data['TARGET']  
X = data.drop(['SK_ID_CURR', 'TARGET'], axis = 1)
```

```

In [3]: experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy", "Tes

def pct(x):
    return round(100*x,1)

class DataFrameSelector(BaseEstimator, TransformerMixin):
    def __init__(self, attribute_names):
        self.attribute_names = attribute_names
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        return X[self.attribute_names].values

def returnModel(x,y,experimentLog,description_text):
    num_attribs = []
    cat_attribs = []

    for col in x.columns.tolist():
        if x[col].dtype in ('int', 'float'):
            num_attribs.append(col)
        else:
            cat_attribs.append(col)

    le_dict = {}
    for col in x.columns.tolist():
        if X[col].dtype == 'object':
            le = LabelEncoder()
            x[col] = x[col].fillna("NULL")
            x[col] = le.fit_transform(x[col])
            le_dict['le_{}'.format(col)] = le

    num_pipeline = Pipeline([('selector', DataFrameSelector(num_attribs)),
                             ('scaler', StandardScaler()),
                             ('imputer', SimpleImputer(strategy = 'median'))
                             ])

    cat_pipeline = Pipeline([
        ('selector', DataFrameSelector(cat_attribs)),
        ('imputer', SimpleImputer(strategy='most_frequent'))
    ])

    full_pipeline = FeatureUnion(transformer_list=[
        ("num_pipeline", num_pipeline),
        ("cat_pipeline", cat_pipeline),
    ])

    np.random.seed(42)
    full_pipeline_with_predictor = Pipeline([
        ("preparation", num_pipeline),
        ("linear", LogisticRegression(random_state=42))
    ])

```

```

# split 20% test data with random seed set to 42 for correct experimentLog
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_state=42)
x_train, x_valid, y_train, y_valid = train_test_split(x_train, y_train, test_size=0.20, random_state=42)
print("train data set: ")
print(x_train.shape, y_train.shape)
print("test data set: ")
print(x_test.shape, y_test.shape)
print("validation data set: ")
print(x_valid.shape, y_valid.shape)

start = time()
full_pipeline_with_predictor.fit(x_train, y_train)
np.random.seed(42)

cv30Splits = ShuffleSplit(n_splits = 30, test_size = 0.3, random_state = 0)
logit_scores = cross_val_score(full_pipeline_with_predictor, x_train, y_train, cv=cv30Splits)
logit_score_train = logit_scores.mean()
train_time = np.round(time() - start, 4)

# Time and score test predictions
start = time()
logit_score_test = full_pipeline_with_predictor.score(x_test, y_test)
test_time = np.round(time() - start, 4)

start = time()
logit_score_valid = full_pipeline_with_predictor.score(x_valid, y_valid)
valid_time = np.round(time() - start, 4)

AUC = roc_auc_score(y_test, full_pipeline_with_predictor.predict(x_test))
print("AUC is {}".format(AUC))
print("\n.....\n")
print("Confusion Matrix: {}".format(confusion_matrix(y_test, full_pipeline_with_predictor.predict(x_test))))

no_of_inputs = x.shape[1]

temp_df = pd.DataFrame()
temp_df = temp_df.append(pd.Series(["Baseline with {} inputs".format(no_of_inputs),
                                   AUC, train_time, test_time, "{} - Baseline LogisticRegression".format(no_of_inputs)]))
temp_df.columns = experimentLog.columns

experimentLog = experimentLog.append(temp_df, ignore_index=True)

return le_dict, full_pipeline_with_predictor, experimentLog

```

```
In [4]: le_dict, full_pipeline_with_predictor, experimentLog = returnModel(X,y,experimentLog)

train data set:
(196806, 120) (196806,)
test data set:
(61503, 120) (61503,)
validation data set:
(49202, 120) (49202,)
AUC is 0.5043329019692199

.....

Confusion Matrix: [[56507    47]
 [ 4902    47]]
```

```
In [5]: experimentLog
```

Out[5]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticRegres...

Additional EDA

Discarding features with null values more than 30%

```
In [6]: null_data = X.isna().sum().reset_index().rename(columns={'index':'col_name',0:'null_count'})
null_data['count_%'] = null_data['null_count']/len(X)*100
null_data = null_data[null_data['count_%'] <= 30]
null_data
```

```
Out[6]:
```

	col_name	null_count	count_%
0	NAME_CONTRACT_TYPE	0	0.000000
1	CODE_GENDER	0	0.000000
2	FLAG_OWN_CAR	0	0.000000
3	FLAG_OWN_REALTY	0	0.000000
4	CNT_CHILDREN	0	0.000000
...
115	AMT_REQ_CREDIT_BUREAU_DAY	41519	13.501631
116	AMT_REQ_CREDIT_BUREAU_WEEK	41519	13.501631
117	AMT_REQ_CREDIT_BUREAU_MON	41519	13.501631
118	AMT_REQ_CREDIT_BUREAU_QRT	41519	13.501631
119	AMT_REQ_CREDIT_BUREAU_YEAR	41519	13.501631

75 rows × 3 columns

```
In [7]: selected_columns = null_data['col_name'].tolist() + ['TARGET']
print(selected_columns)
```

```
['NAME_CONTRACT_TYPE', 'CODE_GENDER', 'FLAG_OWN_CAR', 'FLAG_OWN_REALTY', 'CNT_CHILDREN', 'AMT_INCOME_TOTAL', 'AMT_CREDIT', 'AMT_ANNUITY', 'AMT_GOODS_PRICE', 'NAME_TYPE_SUITE', 'NAME_INCOME_TYPE', 'NAME_EDUCATION_TYPE', 'NAME_FAMILY_STAT', 'NAME_HOUSING_TYPE', 'REGION_POPULATION_RELATIVE', 'DAYS_BIRTH', 'DAYS_EMPLOYED', 'DAYS_REGISTRATION', 'DAYS_ID_PUBLISH', 'FLAG_MOBIL', 'FLAG_EMP_PHONE', 'FLAG_WORK_PHONE', 'FLAG_CONT_MOBILE', 'FLAG_PHONE', 'FLAG_EMAIL', 'OCCUPATION_TYPE', 'CNT_FAM_MEMBERS', 'REGION_RATING_CLIENT', 'REGION_RATING_CLIENT_W_CITY', 'WEEKDAY_APPR_PROCESS_START', 'HOUR_APPR_PROCESS_START', 'REG_REGION_NOT_LIVE_REGION', 'REG_REGION_NOT_WORK_REGION', 'LIVE_REGION_NOT_WORK_REGION', 'REG_CITY_NOT_LIVE_CITY', 'REG_CITY_NOT_WORK_CITY', 'LIVE_CITY_NOT_WORK_CITY', 'ORGANIZATION_TYPE', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'FONDKAPREMONT_MODE', 'HOUSETYPE_MODE', 'WALLSMATERIAL_MODE', 'EMERGENCYSTATE_MODE', 'OBS_30_CNT_SOCIAL_CIRCLE', 'DEF_30_CNT_SOCIAL_CIRCLE', 'OBS_60_CNT_SOCIAL_CIRCLE', 'DEF_60_CNT_SOCIAL_CIRCLE', 'DAYS_LAST_PHONE_CHANGE', 'FLAG_DOCUMENT_2', 'FLAG_DOCUMENT_3', 'FLAG_DOCUMENT_4', 'FLAG_DOCUMENT_5', 'FLAG_DOCUMENT_6', 'FLAG_DOCUMENT_7', 'FLAG_DOCUMENT_8', 'FLAG_DOCUMENT_9', 'FLAG_DOCUMENT_10', 'FLAG_DOCUMENT_11', 'FLAG_DOCUMENT_12', 'FLAG_DOCUMENT_13', 'FLAG_DOCUMENT_14', 'FLAG_DOCUMENT_15', 'FLAG_DOCUMENT_16', 'FLAG_DOCUMENT_17', 'FLAG_DOCUMENT_18', 'FLAG_DOCUMENT_19', 'FLAG_DOCUMENT_20', 'FLAG_DOCUMENT_21', 'AMT_REQ_CREDIT_BUREAU_HOUR', 'AMT_REQ_CREDIT_BUREAU_DAY', 'AMT_REQ_CREDIT_BUREAU_WEEK', 'AMT_REQ_CREDIT_BUREAU_MON', 'AMT_REQ_CREDIT_BUREAU_QRT', 'AMT_REQ_CREDIT_BUREAU_YEAR', 'TARGET']
```

```
In [8]: null_data['col_type'] = null_data['col_name'].apply(lambda x: X[x].dtype)
null_data[null_data['count_%'] > 0]
```

```
Out[8]:
```

	col_name	null_count	count_%	col_type
7	AMT_ANNUITY	12	0.003902	float64
8	AMT_GOODS_PRICE	278	0.090403	float64
27	CNT_FAM_MEMBERS	2	0.000650	float64
40	EXT_SOURCE_2	660	0.214626	float64
41	EXT_SOURCE_3	60965	19.825307	float64
89	OBS_30_CNT_SOCIAL_CIRCLE	1021	0.332021	float64
90	DEF_30_CNT_SOCIAL_CIRCLE	1021	0.332021	float64
91	OBS_60_CNT_SOCIAL_CIRCLE	1021	0.332021	float64
92	DEF_60_CNT_SOCIAL_CIRCLE	1021	0.332021	float64
93	DAYS_LAST_PHONE_CHANGE	1	0.000325	float64
114	AMT_REQ_CREDIT_BUREAU_HOUR	41519	13.501631	float64
115	AMT_REQ_CREDIT_BUREAU_DAY	41519	13.501631	float64
116	AMT_REQ_CREDIT_BUREAU_WEEK	41519	13.501631	float64
117	AMT_REQ_CREDIT_BUREAU_MON	41519	13.501631	float64
118	AMT_REQ_CREDIT_BUREAU_QRT	41519	13.501631	float64
119	AMT_REQ_CREDIT_BUREAU_YEAR	41519	13.501631	float64


Filling null values in NAME_TYPE_SUITE column with "Other_C"

```
In [9]: X_feature = data[selected_columns]
X_feature['NAME_TYPE_SUITE'].fillna('Other_C', inplace=True)
X_feature.head()
```

```
Out[9]:
```

	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN
0	Cash loans	M	N	Y	
1	Cash loans	F	N	N	
2	Revolving loans	M	Y	Y	
3	Cash loans	F	N	Y	
4	Cash loans	M	N	Y	

5 rows × 76 columns



Filling null values in the columns containing keyword AMT_REQ_CREDIT column with 0


```
In [10]: temp_col_reqd = null_data[null_data['null_count'] != 0].reset_index(drop=True)
for col in temp_col_reqd:
    if 'AMT_REQ_CREDIT' in col:
        print("columns to be filled with 0 is: {}".format(col))
        X_feature[col].fillna(0,inplace=True)
```

columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_HOUR
columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_DAY
columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_WEEK
columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_MON
columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_QRT
columns to be filled with 0 is: AMT_REQ_CREDIT_BUREAU_YEAR

Filling null values in the column containing keyword CNT_SOCIAL_CIRCLE with 0

```
In [11]: for col in temp_col_reqd:
    if 'CNT_SOCIAL_CIRCLE' in col:
        print("columns to be filled with 0 is: {}".format(col))
        X_feature[col].fillna(0,inplace=True)
```

columns to be filled with 0 is: OBS_30_CNT_SOCIAL_CIRCLE
columns to be filled with 0 is: DEF_30_CNT_SOCIAL_CIRCLE
columns to be filled with 0 is: OBS_60_CNT_SOCIAL_CIRCLE
columns to be filled with 0 is: DEF_60_CNT_SOCIAL_CIRCLE

Filling null values in the column CNT_FAM_MEMBERS with median

```
In [12]: for col in temp_col_reqd:
    if 'CNT_FAM_MEMBERS' in col:
        print("columns to be filled with median is: {}".format(col))
        X_feature[col].fillna(X_feature[col].median(),inplace=True)
```

columns to be filled with median is: CNT_FAM_MEMBERS


Filling null values in the column AMT_GOODS_PRICE with median for the respective category

```
In [13]: temp_plt_data = X_feature[['AMT_GOODS_PRICE', 'NAME_FAMILY_STATUS']]
temp_plt_data = temp_plt_data.groupby('NAME_FAMILY_STATUS')['AMT_GOODS_PRICE'].mean()
temp_plt_data['AMT_GOODS_PRICE'] = temp_plt_data['AMT_GOODS_PRICE'].fillna(temp_plt_data.head())
```


```
Out[13]:
```

	NAME_FAMILY_STATUS	AMT_GOODS_PRICE
0	Civil marriage	450000.0
1	Married	459000.0
2	Separated	450000.0
3	Single / not married	373500.0
4	Unknown	360000.0

```
In [14]: sns.set(rc={'figure.figsize':(11.7,8.27)})
ax = sns.barplot(x="NAME_FAMILY_STATUS", y="AMT_GOODS_PRICE", data=temp_plt_data)
```



```
In [15]: sns.set(rc={'figure.figsize':(11.7,8.27)})  
ax = sns.boxplot(x="NAME_FAMILY_STATUS", y="AMT_GOODS_PRICE", data=X_feature)
```



```
In [16]: def fill_category_value(a):
    if a['AMT_GOODS_PRICE'] == np.inf:
        return temp_plt_data[temp_plt_data['NAME_FAMILY_STATUS']==a['NAME_FAMILY_STATUS']]['AMT_GOODS_PRICE'].median()
    else:
        return a['AMT_GOODS_PRICE']

    for col in temp_col_reqd:
        X_feature['AMT_GOODS_PRICE'] = X_feature['AMT_GOODS_PRICE'].fillna(np.inf)
        if 'AMT_GOODS_PRICE' in col:
            print("columns to be filled with category median is: {}".format(col))
            X_feature['AMT_GOODS_PRICE'] = X_feature.apply(lambda a: fill_category_value(a), axis=1)
```

columns to be filled with category median is: AMT_GOODS_PRICE

Dropping one single row with column DAYS_LAST_PHONE_CHANGE as null

```
In [17]: X_feature.dropna(subset=['DAYS_LAST_PHONE_CHANGE'], inplace=True)
```

Dropping 12 rows with column AMT_ANNUITY as null

```
In [18]: X_feature.dropna(subset=['AMT_ANNUITY'], inplace=True)
```

```
In [19]: X_feature = X_feature.reset_index(drop=True)
```

Checking highest correlated features with External Source to replace the null values with

Check for EXT_SOURCE_2

```
In [20]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

temp_corr_df = pd.DataFrame()

for col in X_feature.columns.tolist():
    if X_feature[col].dtype == 'int':
        l = [col, X_feature['EXT_SOURCE_2'].corr(X_feature[col])]
    else:
        l = [col, X_feature['EXT_SOURCE_2'].corr(pd.DataFrame(LabelEncoder().fit(X_feature['EXT_SOURCE_2']).transform(X_feature['EXT_SOURCE_2']).values).values)]
    temp_corr_df = temp_corr_df.append(pd.Series(l), ignore_index=True)
temp_corr_df = temp_corr_df.rename(columns={0: 'col_name', 1: 'correlation_with_EXT_2'})
temp_corr_df['correlation_with_EXT_2'] = abs(temp_corr_df['correlation_with_EXT_2'])
temp_corr_df.sort_values(by='correlation_with_EXT_2', ascending=False).head(6).
```

```
Out[20]:
```

	col_name	correlation_with_EXT_2
27	REGION_RATING_CLIENT	0.292903
28	REGION_RATING_CLIENT_W_CITY	0.288306
48	DAYS_LAST_PHONE_CHANGE	0.195766
5	AMT_INCOME_TOTAL	0.170547
75	TARGET	0.160471

REGION_RATING_CLIENT : Our rating of the region where our client lives

```
In [21]: region_rating_grouped = X_feature.groupby('REGION_RATING_CLIENT')['EXT_SOURCE_2']

def fill_external_source2(a):
    if a['EXT_SOURCE_2'] == np.inf:
        return region_rating_grouped[region_rating_grouped['REGION_RATING_CLIENT'] == a['REGION_RATING_CLIENT']]['EXT_SOURCE_2'].mean()
    else:
        return a['EXT_SOURCE_2']

X_feature['EXT_SOURCE_2'] = X_feature['EXT_SOURCE_2'].fillna(np.inf)
X_feature['EXT_SOURCE_2'] = X_feature.apply(lambda a: fill_external_source2(a), axis=1)
```

Check for EXT_SOURCE_3

```
In [22]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

temp_corr_df = pd.DataFrame()

for col in X_feature.columns.tolist():
    if X_feature[col].dtype == 'int':
        l = [col, X_feature['EXT_SOURCE_3'].corr(X_feature[col])]
    else:
        l = [col, X_feature['EXT_SOURCE_3'].corr(pd.DataFrame(LabelEncoder().fit(
temp_corr_df = temp_corr_df.append(pd.Series(l), ignore_index=True)
temp_corr_df = temp_corr_df.rename(columns={0: 'col_name', 1: 'correlation_with_EXT_
temp_corr_df['correlation_with_EXT_3'] = abs(temp_corr_df['correlation_with_EXT_
temp_corr_df = temp_corr_df.sort_values(by='correlation_with_EXT_3', ascending=False)
temp_corr_df
```

```
Out[22]:
```

	col_name	correlation_with_EXT_3
15	DAYS_BIRTH	0.205474
75	TARGET	0.178929
18	DAYS_ID_PUBLISH	0.131598
20	FLAG_EMP_PHONE	0.115284
16	DAYS_EMPLOYED	0.113426
38	EXT_SOURCE_2	0.109728
17	DAYS_REGISTRATION	0.107570
5	AMT_INCOME_TOTAL	0.088906
37	ORGANIZATION_TYPE	0.087994

```
In [23]: ext_source_data = X_feature[temp_corr_df['col_name'].tolist()+['EXT_SOURCE_3']]

for col in ext_source_data.columns.tolist():
    if col != 'EXT_SOURCE_3':
        ext_source_data[col] = LabelEncoder().fit_transform(X_feature[[col]])

ext_source3_train = ext_source_data[ext_source_data['EXT_SOURCE_3'].notnull()]
ext_source3_test = ext_source_data[ext_source_data['EXT_SOURCE_3'].isnull()]
ext_source3_train.shape, ext_source3_test.shape
```

```
Out[23]: ((246535, 10), (60963, 10))
```

```
In [24]: exs3_y_train = ext_source3_train[['EXT_SOURCE_3']]
exs3_X_train = ext_source3_train.drop(columns=['EXT_SOURCE_3'])

exs3_X_test = ext_source3_test.drop(columns=['EXT_SOURCE_3'])
```

```
In [25]: from sklearn.linear_model import LinearRegression

model = LinearRegression().fit(exs3_X_train, exs3_y_train)
y_pred_exs3 = model.predict(exs3_X_test)

exs3_output = exs3_X_test
exs3_output['exs3_y'] = y_pred_exs3
exs3_output
```

```
Out[25]:
```

	DAYS_BIRTH	TARGET	DAYS_ID_PUBLISH	FLAG_EMP_PHONE	DAYS_EMPLOYED	EXT_SOURCE_3
1	8382	0	5876	1	11384	
3	6142	0	3730	1	9533	
4	5215	0	2709	1	9534	
9	10676	0	2175	1	10553	
14	10562	0	4111	1	12369	
...
307471	12298	0	6132	1	12244	
307488	12184	0	2387	1	11526	
307491	8442	0	5908	1	5318	
307493	15818	0	4185	1	12336	
307494	4372	0	2077	0	12573	

60963 rows × 10 columns

```
In [26]: exs3_output = exs3_output.reset_index().rename(columns={'index':'index_to_be_updated'})
for i in exs3_output['index_to_be_updated'].tolist():
    X_feature['EXT_SOURCE_3'].iloc[i] = exs3_output[exs3_output['index_to_be_updated'] == i]['EXT_SOURCE_3'].iloc[0]
```

Checking the null values in the Train dataset

```
In [27]: new_null_data = X_feature.isna().sum().reset_index().rename(columns={'index':'col_name'})
new_null_data['count_%'] = new_null_data['null_count']/len(X_feature)*100
new_null_data = new_null_data[new_null_data['count_%'] <= 30]
new_null_data['col_type'] = new_null_data['col_name'].apply(lambda x: X_feature[x].dtype)
new_null_data[new_null_data['count_%'] > 0]
```

```
Out[27]:
```

col_name	null_count	count_%	col_type
----------	------------	---------	----------

```
In [28]: X_feature.isna().sum()
```

```
Out[28]: NAME_CONTRACT_TYPE      0
        CODE_GENDER              0
        FLAG_OWN_CAR              0
        FLAG_OWN_REALTY           0
        CNT_CHILDREN              0
        ..
        AMT_REQ_CREDIT_BUREAU_WEEK 0
        AMT_REQ_CREDIT_BUREAU_MON  0
        AMT_REQ_CREDIT_BUREAU_QRT  0
        AMT_REQ_CREDIT_BUREAU_YEAR 0
        TARGET                    0
        Length: 76, dtype: int64
```

Training and testing with the selected columns

Adding Additional relevant features felt

```
In [29]: X_feature['AMT_CREDIT_TO_ANNUIITY_RATIO'] = X_feature['AMT_CREDIT'] / X_feature['AMT_INCOME_TOTAL']
        X_feature['Tot_EXTERNAL_SOURCE'] = X_feature['EXT_SOURCE_2'] + X_feature['EXT_SOURCE_1']
        X_feature['Salary_to_credit'] = X_feature['AMT_INCOME_TOTAL']/X_feature['AMT_CREDIT']
        X_feature['Annuity_to_salary_ratio'] = X_feature['AMT_ANNUIITY']/X_feature['AMT_INCOME_TOTAL']
```

```
In [30]: y = X_feature[['TARGET']]
        X = X_feature.drop(['TARGET'], axis = 1)
```

```
In [31]: le_dict, full_pipeline_with_predictor, experimentLog = returnModel(X,y,experimentLog)
```

```
train data set:
(196798, 79) (196798, 1)
test data set:
(61500, 79) (61500, 1)
validation data set:
(49200, 79) (49200, 1)
AUC is 0.5055202479077585
```

.....

```
Confusion Matrix: [[56482    60]
 [ 4898    60]]
```


In [32]: experimentLog

Out[32]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticReges...
1	Baseline with 79 inputs	91.9	91.9	91.8	0.505520	96.4481	0.0506	Selected features Dataset - Baseline LogisticR...

Hyperparameter tuning(Grid Search)

In [33]: train = X_feature
train.shape

Out[33]: (307498, 80)

In [34]: features = train.columns.tolist()
features.remove('TARGET')
len(features)

Out[34]: 79

In [35]: le_dict = {}

for col in features:
 if train[col].dtype == 'object':
 le = LabelEncoder()
 train[col] = le.fit_transform(train[col])
 le_dict['le_{}'.format(col)] = le

In [36]: X = train[features]
y = train['TARGET']

In [37]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state=42)
X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size=0.15, random_state=42)

Decision Making Tree

```

In [38]: DMT_pipe = Pipeline([
            ('scaler', StandardScaler()),
            ('regressor', DecisionTreeRegressor(random_state=100))
        ])

param_grid = [{
    'regressor__max_depth': [2, 3],
    'regressor__min_samples_split': [2, 3],
}]

start = time()
dmt_search = GridSearchCV(estimator=DMT_pipe, param_grid=param_grid, cv=2)
dmt_search.fit(X_train, y_train)
train_time = np.round(time() - start, 4)

trainAcc = dmt_search.score(X_train, y_train)
validAcc = dmt_search.score(X_valid, y_valid)
start = time()
testAcc = dmt_search.score(X_test, y_test)
test_time = np.round(time() - start, 4)

number_of_inputs = X_train.shape[1]
AUC = roc_auc_score(y_test, dmt_search.predict(X_test))

try: experimentLog
except : experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy",
                                                "AUC", "Train Time(s)", "Test Time(s)",
                                                "Experiment description"])
experimentLog.loc[len(experimentLog)] = [f"Gridsearch Decision Making Tree with {number_of_inputs} inputs",
                                          f"{trainAcc*100:8.2f}%", f"{testAcc*100:8.2f}%",
                                          train_time, test_time,
                                          "Decision Making Tree GridSearch with selected features"]

```

In [39]: experimentLog

Out[39]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticRegres...
1	Baseline with 79 inputs	91.9	91.9	91.8	0.505520	96.4481	0.0506	Selected features Dataset - Baseline LogisticR...
2	Gridsearch Decision Making Tree with 79 inputs	7.84%	7.56%	7.35%	0.738725	4.4617	0.0129	Decision Making Tree GridSearch with selected ...

Lasso Regression

```

In [40]: lasso_pipeline = Pipeline([
        ('scaler', StandardScaler()),
        ('model', Lasso())
    ])

lasso_pipeline = GridSearchCV(lasso_pipeline,
    {'model__alpha': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]},
    cv = 5, scoring="neg_mean_squared_error", verbose=3
)

start = time()
lasso_pipeline.fit(X_train, y_train)
train_time = np.round(time() - start, 4)

trainAcc = lasso_pipeline.score(X_train, y_train)
validAcc = lasso_pipeline.score(X_valid, y_valid)
start = time()
testAcc = lasso_pipeline.score(X_test, y_test)
test_time = np.round(time() - start, 4)

AUC = roc_auc_score(y_test, lasso_pipeline.predict(X_test))
number_of_inputs = X_train.shape[1]

#del experimentLog
try: experimentLog
except : experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy", "AUC", "Train Time(s)", "Test Time(s)"])
experimentLog.loc[len(experimentLog)] = [f"Lasso Reg with {number_of_inputs} inputs",
    f"{trainAcc*100:8.2f}%", f"{testAcc*100:8.2f}%",
    train_time, test_time,
    "Lasso Regression for feature selection"]

```

Fitting 5 folds for each of 22 candidates, totalling 110 fits

```

[CV 1/5] END .....model__alpha=0.0001; score=-0.070 total time= 1
5.2s
[CV 2/5] END .....model__alpha=0.0001; score=-0.070 total time= 1
5.5s
[CV 3/5] END .....model__alpha=0.0001; score=-0.067 total time= 1
5.7s
[CV 4/5] END .....model__alpha=0.0001; score=-0.069 total time= 1
5.5s
[CV 5/5] END .....model__alpha=0.0001; score=-0.069 total time= 1
5.3s
[CV 1/5] END .....model__alpha=0.001; score=-0.070 total time=
2.1s
[CV 2/5] END .....model__alpha=0.001; score=-0.070 total time=
2.2s
[CV 3/5] END .....model__alpha=0.001; score=-0.067 total time=
2.2s
[CV 4/5] END .....model__alpha=0.001; score=-0.070 total time=
2.2s
[CV 5/5] END .....model__alpha=0.001; score=-0.069 total time=
2.1s

```

In [41]: experimentLog

Out[41]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticRegres...
1	Baseline with 79 inputs	91.9	91.9	91.8	0.505520	96.4481	0.0506	Selected features Dataset - Baseline LogisticR...
2	Gridsearch Decision Making Tree with 79 inputs	7.84%	7.56%	7.35%	0.738725	4.4617	0.0129	Decision Making Tree GridSearch with selected ...
3	Lasso Reg with 79 inputs	-6.89%	-6.87%	-6.89%	0.755827	131.6072	0.0224	Lasso Regression for feature selection

```
In [42]: print(lasso_pipeline.best_params_)
coefficients = lasso_pipeline.best_estimator_.named_steps['model'].coef_
importance = np.abs(coefficients)
len(np.array(features)[importance > 0])

{'model__alpha': 0.0001}
```

Out[42]: 74

Ridge Regression

```
In [43]: ridge_pipeline = Pipeline([
        ('scaler', StandardScaler()),
        ('model', Ridge())
    ])

ridge_pipeline = GridSearchCV(ridge_pipeline,
                              {'model__alpha': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]},
                              cv = 5, scoring="neg_mean_squared_error", verbose=3
                              )

start = time()
ridge_pipeline.fit(X_train, y_train)
train_time = np.round(time() - start, 4)

trainAcc = ridge_pipeline.score(X_train, y_train)
validAcc = ridge_pipeline.score(X_valid, y_valid)
start = time()
testAcc = ridge_pipeline.score(X_test, y_test)
test_time = np.round(time() - start, 4)

AUC = roc_auc_score(y_test, ridge_pipeline.predict(X_test))
number_of_inputs = X_train.shape[1]

try: experimentLog
except : experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy", "AUC", "Train Time(s)", "Test Time(s)"])
experimentLog.loc[len(experimentLog)] = [f"Ridge Reg with {number_of_inputs} inputs",
                                         f"{trainAcc*100:8.2f}%", f"{testAcc*100:8.2f}%",
                                         train_time, test_time,
                                         "Ridge Regression for feature selection"]
```

Fitting 5 folds for each of 22 candidates, totalling 110 fits

```
[CV 1/5] END .....model__alpha=0.0001; score=-0.070 total time=0.3s
[CV 2/5] END .....model__alpha=0.0001; score=-0.070 total time=0.3s
[CV 3/5] END .....model__alpha=0.0001; score=-0.067 total time=0.3s
[CV 4/5] END .....model__alpha=0.0001; score=-0.069 total time=0.3s
[CV 5/5] END .....model__alpha=0.0001; score=-0.069 total time=0.3s
[CV 1/5] END .....model__alpha=0.001; score=-0.070 total time=0.3s
[CV 2/5] END .....model__alpha=0.001; score=-0.070 total time=0.3s
[CV 3/5] END .....model__alpha=0.001; score=-0.067 total time=0.2s
[CV 4/5] END .....model__alpha=0.001; score=-0.069 total time=0.3s
[CV 5/5] END .....model__alpha=0.001; score=-0.069 total time=0.3s
```

In [44]: experimentLog

Out[44]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticRegres...
1	Baseline with 79 inputs	91.9	91.9	91.8	0.505520	96.4481	0.0506	Selected features Dataset - Baseline LogisticR...
2	Gridsearch Decision Making Tree with 79 inputs	7.84%	7.56%	7.35%	0.738725	4.4617	0.0129	Decision Making Tree GridSearch with selected ...
3	Lasso Reg with 79 inputs	-6.89%	-6.87%	-6.89%	0.755827	131.6072	0.0224	Lasso Regression for feature selection
4	Ridge Reg with 79 inputs	-6.89%	-6.87%	-6.89%	0.756776	30.4069	0.0134	Ridge Regression for feature selection

```
In [45]: print(ridge_pipeline.best_params_)
coefficients = ridge_pipeline.best_estimator_.named_steps['model'].coef_
importance = np.abs(coefficients)
len(np.array(features)[importance > 0.005])

{'model__alpha': 3}
```

Out[45]: 22

Logistic Regression

```

In [46]: logistic = LogisticRegression(max_iter=10000, tol=0.1)
         clf_pipe = Pipeline(steps=[("logistic", logistic)])
         param_grid = {
             "logistic__C": np.logspace(-4, 4, 10)
         }

         # Time and score test predictions
         start = time()

         clf_search = GridSearchCV(clf_pipe, param_grid, n_jobs=-1)
         clf_search.fit(X_train, y_train)
         train_time = np.round(time() - start, 4)

         trainAcc = clf_search.score(X_train, y_train)
         validAcc = clf_search.score(X_valid, y_valid)
         start = time()
         testAcc = clf_search.score(X_test, y_test)
         test_time = np.round(time() - start, 4)

         number_of_inputs = X_train.shape[1]
         AUC = roc_auc_score(y_test, clf_search.predict(X_test))

         try: experimentLog
         except : experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy",
                                                         "AUC", "Train Time(s)", "Test Time(s)"])
         experimentLog.loc[len(experimentLog)] = [f"Gridsearch LogReg with {number_of_inputs} features",
                                                  f"{trainAcc*100:8.2f}%", f"{testAcc*100:8.2f}%",
                                                  train_time, test_time,
                                                  "LogReg GridSearch with selected features"]

```

```
In [47]: experimentLog
```

Out[47]:

	ExpID	Cross fold train accuracy	Test Accuracy	Validation Accuracy	AUC	Train Time(s)	Test Time(s)	Experiment description
0	Baseline with 120 inputs	92.0	92.0	91.8	0.504333	110.2988	0.0994	All features Dataset - Baseline LogisticRegres...
1	Baseline with 79 inputs	91.9	91.9	91.8	0.505520	96.4481	0.0506	Selected features Dataset - Baseline LogisticR...
2	Gridsearch Decision Making Tree with 79 inputs	7.84%	7.56%	7.35%	0.738725	4.4617	0.0129	Decision Making Tree GridSearch with selected ...
3	Lasso Reg with 79 inputs	-6.89%	-6.87%	-6.89%	0.755827	131.6072	0.0224	Lasso Regression for feature selection
4	Ridge Reg with 79 inputs	-6.89%	-6.87%	-6.89%	0.756776	30.4069	0.0134	Ridge Regression for feature selection
5	Gridsearch LogReg with 79 inputs	91.91%	91.98%	91.96%	0.500000	55.1705	0.0168	LogReg GridSearch with selected features

Xgboost


```

In [ ]: xgboost_pipe = Pipeline([
        ('standard_scaler', StandardScaler()),
        ('model', xgb.XGBClassifier())
    ])

param_grid = {
    'model__max_depth': [2, 3, 5, 7, 10],
    'model__n_estimators': [10, 100, 500]
}

xgboost_search = GridSearchCV(xgboost_pipe, param_grid, scoring="roc_auc", cv=5)

# Time and score test predictions
start = time()
xgboost_search.fit(X_train, y_train)
train_time = np.round(time() - start, 4)

trainAcc = xgboost_search.score(X_train, y_train)
validAcc = xgboost_search.score(X_valid, y_valid)
start = time()
testAcc = xgboost_search.score(X_test, y_test)
test_time = np.round(time() - start, 4)

number_of_inputs = X_train.shape[1]
AUC = roc_auc_score(y_test, xgboost_search.predict(X_test))

try: experimentLog
except : experimentLog = pd.DataFrame(columns=["ExpID", "Cross fold train accuracy",
                                                "AUC", "Train Time(s)", "Test Time(s)"])
experimentLog.loc[len(experimentLog)] = [f"Gridsearch Xgboost with {number_of_inputs} inputs",
                                          f"{trainAcc*100:8.2f}%", f"{testAcc*100:8.2f}%",
                                          train_time, test_time,
                                          "Xgboost GridSearch with selected features"]

```

```

In [ ]: experimentLog

```